



S&DS 265 / 565  
Introductory Machine Learning

# Word Embeddings

Thursday, October 28

ADV

ADJ

NOUN

VERB

PRON

Yale

# Reminders

- Assignment 4 due Tuesday
- Assignment 5 out Tuesday

# Language models

- A language model is a way of assigning a probability to any sequence of words (or string of text)

$$p(w_1, \dots, w_n)$$

# Language models

- A language model is a way of assigning a probability to any sequence of words (or string of text)

$$p(w_1, \dots, w_n)$$

- By the basic rules of conditional probability we can factor this as

$$p(w_1, \dots, w_n) = p(w_1)p(w_2 | w_1) \dots p(w_n | w_1, \dots, w_{n-1})$$

# Language models

- A language model is a way of *generating* any sequence of words

$$\begin{aligned} P(\text{"the whole forest had been anesthetized"}) = & \\ & P(\text{"the"}) \times P(\text{"whole"} \mid \text{"the"}) \\ & \times P(\text{"forest"} \mid \text{"the whole"}) \\ & \times P(\text{"had"} \mid \text{"the whole forest"}) \\ & \times P(\text{"been"} \mid \text{"the whole forest had"}) \\ & \times P(\text{"anesthetized"} \mid \text{"the whole forest had been"}) \end{aligned}$$

## Remixing Noon

Text generated from Channel Skin by Jeff Noon

"The whole forest had been anesthetised, her temples wired, her senses stimulated, her eyes were not on Eva, not on Eva, not on Eva, not on anybody in that same realm, the land of dreams and nightmares."

Viability: 0.000000326%



[https://revdancatt.com/2017/03/01/markov\\_noon](https://revdancatt.com/2017/03/01/markov_noon)

# Text generation

- Words generated one-by-one
- A word is chosen by sampling from a probability distribution
- Then treated as if it were “real,” as in dreaming
- Result is purely synthetic text

# How good is a language model? Perplexity

*Perplexity* is defined as

$$\text{Perplexity}(\theta) = \left( \prod_{i=1}^n p_{\theta}(w_i | w_{1:i-1}) \right)^{-\frac{1}{n}}$$

where  $w_1, w_2, \dots, w_n$  is a large chunk of text that wasn't used to train the language model.



# How good is a language model? Perplexity

- Perplexity is the inverse of the geometric mean of the word probabilities
- If the perplexity is 100, the model predicts, on average, as if there were 100 equally likely words to follow
- This is the (geometric) average “branching factor” for the model on real text

# Modern language models

Suppose a computer program assigns a “score” to possible next words:

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

# Modern language models

Suppose a computer program assigns a “score” to possible next words:

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

Can convert this to a language model by the “softmax” operation:

$$p(w \mid w_1, \dots, w_n) = \frac{\exp(s(w; w_1, \dots, w_n))}{\sum_{v \in V} \exp(s(v; w_1, \dots, w_n))}$$

# Modern language models

Suppose a computer program assigns a “score” to possible next words:

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

Can convert this to a language model by the “softmax” operation:

$$p(w \mid w_1, \dots, w_n) = \frac{\exp(s(w; w_1, \dots, w_n))}{\sum_{v \in V} \exp(s(v; w_1, \dots, w_n))}$$

In GPT-3, the function  $s(v; w_{1:n})$  is learned on large amounts of text (unsupervised) using a type of deep neural network called a *transformer*.

# Modern language models

Suppose a computer program assigns a “score” to possible next words:

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

# Modern language models

Suppose a computer program assigns a “score” to possible next words:

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

Today, we'll be working with a simple case where

$$\begin{aligned} s(v; w_1, \dots, w_n) &= \beta_v^T \phi(w_1, \dots, w_n) \\ &= \beta_v^T \phi(w_n) \\ &= \phi(v)^T \phi(w_n) \end{aligned}$$

# Modern language models

Suppose a computer program assigns a “score” to possible next words:

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

Today, we'll be working with a simple case where

$$\begin{aligned} s(v; w_1, \dots, w_n) &= \beta_v^T \phi(w_1, \dots, w_n) \\ &= \beta_v^T \phi(w_n) \\ &= \phi(v)^T \phi(w_n) \end{aligned}$$

# Key intuition

- Words that have similar neighbors will be similar
- Self-referential notion of similarity
- This will be an intuition behind “word embeddings”



# Pointwise mutual information (PMI)

Can cluster words based on “pointwise mutual information” (PMI)

$$\log \left( \frac{p_{\text{near}}(w_1, w_2)}{p(w_1)p(w_2)} \right)$$

- How likely are specific words/clusters to co-occur together within some window, compared to if they were independent?

# Example clusters from PMI

we our us ourselves ours  
question questions asking answer answers answering  
performance performed perform performs performing  
tie jacket suit  
write writes writing written wrote pen  
morning noon evening night nights midnight bed  
attorney counsel trial court judge  
problems problem solution solve analyzed solved solving  
letter addressed enclosed letters correspondence  
large size small larger smaller  
operations operations operating operate operated  
school classroom teaching grade math  
street block avenue corner blocks  
table tables dining chairs plate  
published publication author publish writer titled  
wall ceiling walls enclosure roof  
sell buy selling buying sold

# Core idea of embeddings

- Form a language model but replace classes by vectors, one for each word
- Use PMI-like scores to fit the vectors
- Can be applied whenever have cooccurrence data.

# Constructing embeddings

Language model is

$$p(w_2 | w_1) = \frac{\exp(\phi(w_2)^T \phi(w_1))}{\sum_w \exp(\phi(w)^T \phi(w_1))}.$$

Carry out stochastic gradient descent over the embedding vectors  $\phi \in \mathbb{R}^d$  (where  $d \approx 50$ – $500$  is chosen by hand)

This is what Mikolov et al. (2014, 2015) did at Google. With a couple of twists:

# Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word

# Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word
- This leads to a model of nearby words  $p_{\text{near}}(w_2 | w_1)$ .

# Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word
- This leads to a model of nearby words  $p_{\text{near}}(w_2 | w_1)$ .
- Second is computational. The bottleneck is computing the denominator in the logistic (softmax) probability.

# Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word
- This leads to a model of nearby words  $p_{\text{near}}(w_2 | w_1)$ .
- Second is computational. The bottleneck is computing the denominator in the logistic (softmax) probability.
- Use “negative sampling”: Approximation

$$\begin{aligned} & \sum_w \exp(\phi(w)^T \phi(w_1)) \\ & \approx \exp(\phi(w_2)^T \phi(w_1)) + \sum_{\text{random } w} \exp(\phi(w)^T \phi(w_1)) \end{aligned}$$



# Using PCA

A closely related approach is to use PCA of PMI:

- Form  $V \times V$  matrix of pointwise mutual information values

$$\log \left( \frac{p_{\text{near}}(w_1, w_2)}{p(w_1)p(w_2)} \right)$$

- Compute top  $k$  eigenvectors  $\phi_1, \dots, \phi_k$
- For each word  $w$ , define embedding as

$$\phi(w) \equiv (\phi_{1w}, \phi_{2w}, \dots, \phi_{kw})^T$$

# Analogies

These heuristics enable training on very large text collections. Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

# Analogies

These heuristics enable training on very large text collections. Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

Paris is to France as ? is to Germany

# Analogies

These heuristics enable training on very large text collections. Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

Paris is to France as ? is to Germany

$$\phi(\text{king}) - \phi(\text{man}) \stackrel{?}{\approx} \phi(\text{queen}) - \phi(\text{woman})$$

$$\hat{w} = \arg \min_w \|\phi(\text{king}) - \phi(\text{man}) + \phi(\text{woman}) - \phi(w)\|^2$$

Does  $\hat{w} = \text{queen}$ ?

# Learned Analogies

Table 8: *Examples of the word pair relationships, using the best word vectors from Table 4 (Skip-gram model trained on 783M words with 300 dimensionality).*

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

# Evaluation Analogies

Type of relationship	Word Pair 1		Word Pair 2	
Common capital city	Athens	Greece	Oslo	Norway
All capital cities	Astana	Kazakhstan	Harare	Zimbabwe
Currency	Angola	kwanza	Iran	rial
City-in-state	Chicago	Illinois	Stockton	California
Man-Woman	brother	sister	grandson	granddaughter
Adjective to adverb	apparent	apparently	rapid	rapidly
Opposite	possibly	impossibly	ethical	unethical
Comparative	great	greater	tough	tougher
Superlative	easy	easiest	lucky	luckiest
Present Participle	think	thinking	read	reading
Nationality adjective	Switzerland	Swiss	Cambodia	Cambodian
Past tense	walking	walked	swimming	swam
Plural nouns	mouse	mice	dollar	dollars
Plural verbs	work	works	speak	speaks

# GloVe

Shortly after: Stanford group introduced a variant

$$\mathcal{O}(\phi) = \sum_{w_1, w_2} f(c_{w_1, w_2}) \left( \phi(w_1)^T \phi(w_2) - \log c_{w_1, w_2} \right)^2$$

where  $c_{w, w'}$  are cooccurrence counts in a window (PMI)

- A type of regression estimator
- Main advantage is that SGD can be carried out much more efficiently

# GloVe

$$\mathcal{O}(\phi) = \sum_{w_1, w_2} f(c_{w_1, w_2}) \left( \phi(w_1)^T \phi(w_2) - \log c_{w_1, w_2} \right)^2$$

where  $c_{w, w'}$  are cooccurrence counts.

- Heuristic weighting function

$$f(x) = \left( \frac{x}{x_{\max}} \right)^{\alpha}$$

where  $\alpha = 3/4$  set empirically.

- So  $10^{-4} \mapsto 10^{-3}$ . Each order of magnitude down gets “boosted” by 1/4-magnitude.



# GloVe site and code

## GloVe: Global Vectors for Word Representation

Jeffrey Pennington, Richard Socher, Christopher D. Manning

### Introduction

GloVe is an unsupervised learning algorithm for obtaining vector representations for words. Training is performed on aggregated global word-word co-occurrence statistics from a corpus, and the resulting representations showcase interesting linear substructures of the word vector space.

### Getting started (Code download)

- Download the [code](#) (licensed under the [Apache License, Version 2.0](#))
- Unpack the files: `unzip GloVe-1.2.zip`
- Compile the source: `cd GloVe-1.2 && make`
- Run the demo script, `demo.sh`
- Consult the included README for further usage details, or ask a [question](#)
- The code is also available [on GitHub](#)

### Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](#) v1.0 whose full text can be found at: <http://www.opensource.org/licenses/oddl/1.0/>
  - [WikiPedia 2014 - English word](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download) [glove.6B.zip](#)
  - [Common Crawl](#) (2B tokens, 10M vocab, uncased, 300d vectors, 175 GB download) [glove.42B.300d.zip](#)
  - [Common Crawl](#) (840B tokens, 22M vocab, cased, 300d vectors, 2.03 GB download) [glove.840B.300d.zip](#)
  - [Twitter](#) (2B tweets, 12M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download) [glove.twitter.27B.zip](#)
- Ruby [script](#) for preprocessing Twitter data

### Citing GloVe

Jeffrey Pennington, Richard Socher, and Christopher D. Manning, 2014. [GloVe: Global Vectors for Word Representation](#) ([pdf](#)) ([bib](#))

### Highlights

#### 1. Nearest neighbors

The Euclidean distance (or cosine similarity) between two word vectors provides an effective method for measuring the linguistic or semantic similarity of the corresponding words. Sometimes, the nearest neighbors according to this metric reveal rare but relevant words that lie outside an average human's vocabulary. For example, here are the closest words to the target word *frog*:

1. frog
2. toad
3. *litoria*
4. *leptodactylidae*
5. rana
6. lizard
7. *eleutherodactylus*



3. *litoria*



4. *leptodactylidae*



5. rana





7. *eleutherodactylus*

#### 2. Linear substructures

The similarity metrics used for nearest neighbor evaluations produce a single scalar that quantifies the relatedness of two words. This simplicity can be problematic since two given words almost always exhibit more intricate relationships than can be captured by a single number. For example, men may be regarded as similar to women in that both words describe human beings; on the other hand, the two words are often considered opposites since they highlight a primary axis along which humans differ from one another.

# Recommendation via Embedding



# Notebook

Let's go to the Python notebook!

# Embedding embeddings: t-SNE

- How can we visualize the embeddings?

# Embedding embeddings: t-SNE

- How can we visualize the embeddings?
- We're in a very high dimensional space

# Embedding embeddings: t-SNE

- How can we visualize the embeddings?
- We're in a very high dimensional space
- Could use PCA—this will tend to distort more

# Embedding embeddings: t-SNE

- How can we visualize the embeddings?
- We're in a very high dimensional space
- Could use PCA—this will tend to distort more
- Many visualization techniques exist. A currently popular one is t-SNE: “Student-t Stochastic Neighborhood Embedding”

# t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings



# t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings
- Scale and symmetrize, giving a matrix  $P = [P_{ij}]$

# t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings
- Scale and symmetrize, giving a matrix  $P = [P_{ij}]$
- Represent word  $i$  by  $y_i \in \mathbb{R}^2$ . Use a heavy-tailed distribution (Student-t)

# t-SNE

Here's the idea behind t-SNE:

- Form a language model using the embeddings
- Scale and symmetrize, giving a matrix  $P = [P_{ij}]$
- Represent word  $i$  by  $y_i \in \mathbb{R}^2$ . Use a heavy-tailed distribution (Student-t)
- Select  $y_i$  using stochastic gradient descent

## t-SNE: More info and examples

<https://lvdmaaten.github.io/tsne/>

<http://cs.stanford.edu/people/karpathy/tsnejs/>

Note: This is just a visualization technique, to give intuition for the high dimensional embedding

# Embedding / Visualization Examples

[WebVectors](#)[Similar words](#)[Visualizations](#)[Calculator](#)[2D text](#)[Miscellaneous](#)[Models](#)[About](#)

## WebVectors: word embeddings online

"You shall know a word by the company it keeps." (Firth 1957)


Enter a word to produce a list of its 10 nearest semantic associates.  
English Wikipedia model will be used; for other models, visit [Similar Words](#) tab.

### Semantic associates for **platypus** (computed on [English Wikipedia](#))

Word frequency

☒ High ☒ Medium ☐ Low

1. marsupial 0.642
2. crocodile 0.605
3. kangaroo 0.595
4. turtle 0.595
5. iguana 0.589
6. frog 0.573
7. penguin 0.572
8. wallaby 0.570
9. alligator 0.569
10. opossum 0.568



Similarity threshold ☐ Show tags

• We show only the associates of the same part of speech as your query. All associates can be found at the [Similar Words](#) tab.

<http://vectors.nlp1.eu/explore/embeddings/en/>

# Summary: Word embeddings

- Word embeddings are vector representations of words, learned from cooccurrence statistics
- The models can be viewed in terms of language modeling, pointwise mutual information, and regression
- Surprising semantic relations are encoded in linear relations
- Embeddings improve with more data
- t-SNE is an algorithm for visualizing embeddings

extra slides (optional)



## t-SNE: Detailed algorithm

For each word  $w_i$  compute a language model

$$P_{j|i} \propto \exp \left( -\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)$$

That is:

$$P_{j|i} = \frac{\exp \left( -\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)}{\sum_k \exp \left( -\frac{\|\phi(w_i) - \phi(w_k)\|^2}{2h_i^2} \right)}$$





## t-SNE: Detailed algorithm

For each word  $w_i$  compute a language model

$$P_{j|i} \propto \exp \left( -\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)$$

That is:

$$P_{j|i} = \frac{\exp \left( -\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)}{\sum_k \exp \left( -\frac{\|\phi(w_i) - \phi(w_k)\|^2}{2h_i^2} \right)}$$

Choose the bandwidth  $h_i$  so that the perplexity is, say, 10. This puts the probabilities all on the same scale.



## t-SNE: Detailed algorithm

For each word  $w_i$  compute a language model

$$P_{j|i} \propto \exp \left( -\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)$$



## t-SNE: Detailed algorithm

For each word  $w_i$  compute a language model

$$P_{j|i} \propto \exp \left( -\frac{\|\phi(w_i) - \phi(w_j)\|^2}{2h_i^2} \right)$$

Now form

$$P_{ij} = \frac{1}{2} (P_{j|i} + P_{i|j})$$

as a simple way of symmetrizing.



## t-SNE: Detailed algorithm

Now form Student-t distribution depending on the visualization vectors  $y_i \in \mathbb{R}^2$ :

$$Q_{ij} \propto \left(1 + \|y_i - y_j\|^2\right)^{-1}$$



## t-SNE: Detailed algorithm

Now form Student-t distribution depending on the visualization vectors  $y_i \in \mathbb{R}^2$ :

$$Q_{ij} \propto \left(1 + \|y_i - y_j\|^2\right)^{-1}$$

That is:

$$Q_{ij} = \frac{\left(1 + \|y_i - y_j\|^2\right)^{-1}}{\sum_{k \neq \ell} \left(1 + \|y_k - y_\ell\|^2\right)^{-1}}$$

This has fatter tails than a Gaussian



## t-SNE: Detailed algorithm

Finally, run stochastic gradient descent (SGD) over the vectors  $y_i$  to optimize:

$$\begin{aligned}\hat{y} &= \arg \min \sum_{ij} P_{ij} \log P_{ij} / Q_{ij} \\ &= \arg \max \sum_{ij} P_{ij} \log Q_{ij}\end{aligned}$$

Interpretation: if  $\phi(w_i)$  is very close to  $\phi(w_j)$  then  $y_i$  will be close to  $y_j$ .  
(long distances may be stretched further...)